Multilevel Multilingual

Multilevel Models in Stata, R and Julia

Andrew Grogan-Kaylor

2023-12-01

Table of contents

List of Figures

List of Tables

4

1 Multilevel Models in Stata, R and Julia

1.1 Introduction

Below, I describe the use of Stata, R, and Julia to estimate multilevel models. Because this document is built by Quarto, I describe calling these programs from within a Quarto environment. However, each piece of software could be used individually and separately.

1.2 The Data

The examples below use the `simulated_multilevel_data.dta` file from *Multilevel Thinking*. Here is a direct link to download the data.

Table 1.1: Sample of Simulated Multilevel Data

country	HDI	family	id	group	physical_punishment	warmth	outcome
1	69	1	1.1	2	2	3	59.18
1	69	2	1.2	2	4	0	61.54
1	69	3	1.3	1	4	4	51.87
1	69	4	1.4	2	0	6	51.71
1	69	5	1.5	2	3	2	55.88
1	69	6	1.6	1	5	3	60.78

1.3 The Equation

$$\text{outcome}_{ij} = \beta_0 + \beta_1 \text{warmth}_{ij} + \beta_2 \text{physical punishment}_{ij} + \beta_3 \text{group}_{ij} + \beta_4 \text{HDI}_{ij} + u_{0j} + u_{1j} \times \text{warmth}_{ij} + e_{ij}$$

$$(1.1)$$

1.4 Setup

1.4.1 Stata

I need to use the library `Statamarkdown` to call Stata, or I could run Stata on its own

```
library(Statamarkdown)
```

1.4.2 R

In R, I use the library `lme4` to run multilevel models.

```
library(lme4)
```

1.4.3 Julia

I need to call Julia from R.

```
library(JuliaCall)

julia_setup(JULIA_HOME = "/Applications/Julia-1.8.app/Contents/Resources/julia/bin")
```

1.5 Get Data & Run Models

To explain statistical syntax for each software, I consider the more general case of a multilevel model with dependent variable y, independent variables x and z, clustering variable `group`, and a random slope for x. i is the index for the person, while j is the index for the `group`.

$$y = \beta_0 + \beta_1 x_{ij} + \beta_2 z_{ij} + u_{0j} + u_{1j} \times x_{ij} + e_{ij} \tag{1.2}$$

1.5.1 Stata

In Stata `mixed`, the syntax for a multilevel model of the form described in Equation 1.2 is:

```
mixed y x || group: x
```

1.5.1.1 Get The Data

> 💡 Tip For Running Stata From Quarto
>
> Because I am running Stata from inside a Quarto document, and running Stata in multiple chunks, I need to use the `collectcode=TRUE` option in the first Stata chunk. i.e. my Quarto chunk needs to begin with "`{stata, collectcode=TRUE}`
> See Doug Hemken's excellent documentation on `Statamarkdown` here.

```stata
use simulated_multilevel_data.dta
```

1.5.1.2 Graph

```stata
twoway scatter outcome warmth, xtitle("warmth") ytitle("outcome") title("Outcome by Parent

quietly graph export scatter.png, replace
```

```
> arental Warmth")
```

1.5.1.3 Run The Model

```stata
mixed outcome warmth physical_punishment group HDI || country: warmth
```

```
Performing EM optimization ...

Performing gradient-based optimization:
Iteration 0:   Log likelihood =  -9668.198
Iteration 1:   Log likelihood = -9667.9551
Iteration 2:   Log likelihood = -9667.9534
Iteration 3:   Log likelihood = -9667.9533
Iteration 4:   Log likelihood = -9667.9532

Computing standard errors ...

Mixed-effects ML regression                     Number of obs    =  3,000
Group variable: country                         Number of groups =     30
```

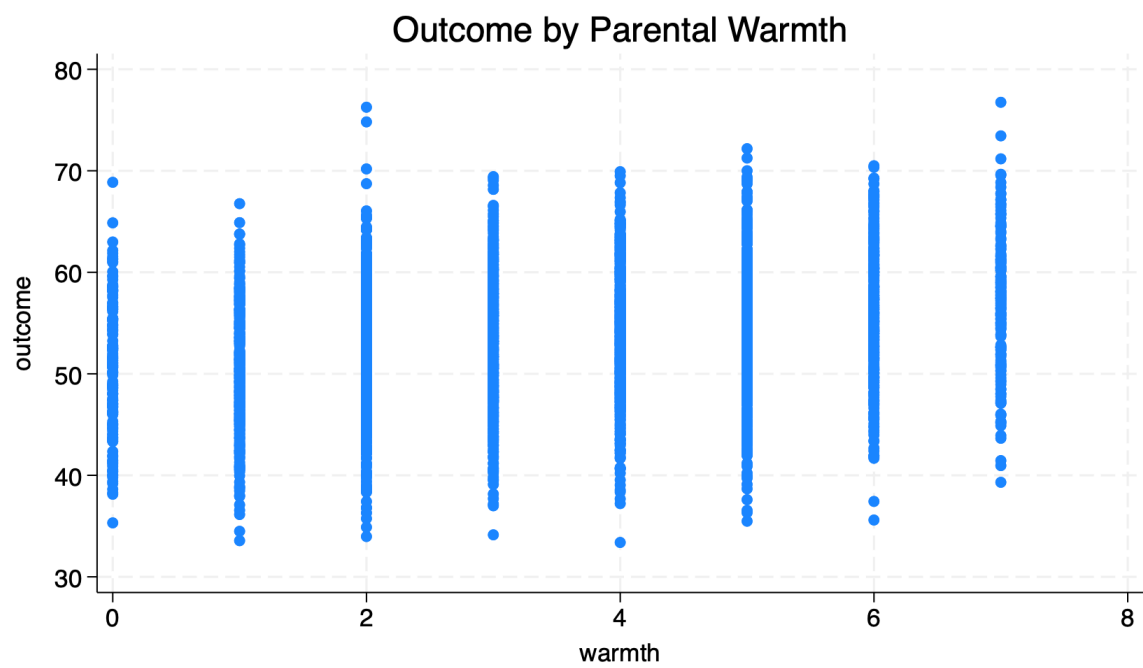Figure 1.1: Outcome by Parental Warmth (Stata)

```
                                              Obs per group:
                                                            min =        100
                                                            avg =      100.0
                                                            max =        100
                                              Wald chi2(4)        = 401.26
Log likelihood = -9667.9532                   Prob > chi2         = 0.0000


--------------------------------------------------------------------------------
          outcome | Coefficient  Std. err.      z    P>|z|     [95% conf. interval]
------------------+-------------------------------------------------------------
           warmth |   .9616447   .0581825    16.53   0.000     .8476091    1.07568
physical_punishment | -.8453802   .0798155   -10.59   0.000    -1.001816   -.6889448
            group |   1.084344   .2200539     4.93   0.000     .6530461   1.515642
              HDI |    .010557   .0204522     0.52   0.606    -.0295286    .0506426
            _cons |   49.87963   1.436612    34.72   0.000     47.06392   52.69534
--------------------------------------------------------------------------------


------------------------------------------------------------------------------
  Random-effects parameters  |   Estimate   Std. err.      [95% conf. interval]
-----------------------------+------------------------------------------------
country: Independent         |
              var(warmth) |   1.83e-06   .0000173     1.76e-14   190.9774
               var(_cons) |   3.370262   .9633726     1.924651   5.901676
-----------------------------+------------------------------------------------
            var(Residual) |   36.01906   .9346936     34.23291   37.89842
------------------------------------------------------------------------------
LR test vs. linear model: chi2(2) = 198.01              Prob > chi2 = 0.0000

Note: LR test is conservative and provided only for reference.
```

1.5.2 R

In R `lme4`, the general syntax for a multilevel model of the form described in Equation 1.2 is:

```
lmer(y ~ x + z + (1 + x || group), data = ...)
```

1.5.2.1 Get The Data

```r
library(haven)

df <- read_dta("simulated_multilevel_data.dta")
```

1.5.2.2 Graph

```r
library(ggplot2)

ggplot(df,
       aes(x = warmth,
           y = outcome)) +
  geom_point() +
  labs(title = "Outcome by Parental Warmth")
```
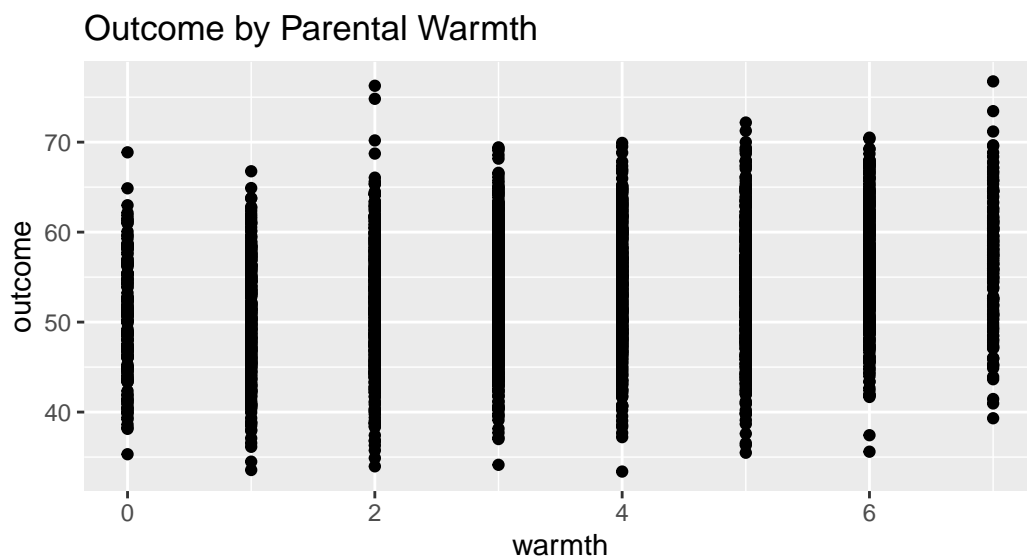
Figure 1.2: Outcome by Parental Warmth (R)

1.5.2.3 Run The Model

```r
fit1 <- lmer(outcome ~ warmth + physical_punishment +
                group + HDI +
                (1 + warmth || country),
             data = df)

summary(fit1)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: outcome ~ warmth + physical_punishment + group + HDI + ((1 |
    country) + (0 + warmth | country))
   Data: df

REML criterion at convergence: 19350.3

Scaled residuals:
    Min      1Q  Median      3Q     Max
-3.4496 -0.6807  0.0016  0.6864  3.1792

Random effects:
 Groups     Name        Variance  Std.Dev.
 country    (Intercept)  3.611568 1.90041
 country.1  warmth       0.001876 0.04331
 Residual               36.049124 6.00409
Number of obs: 3000, groups:  country, 30

Fixed effects:
                     Estimate Std. Error t value
(Intercept)          49.88754    1.48203  33.662
warmth                0.96155    0.05875  16.367
physical_punishment -0.84556    0.07986 -10.588
group                 1.08471    0.22017   4.927
HDI                   0.01044    0.02116   0.493

Correlation of Fixed Effects:
           (Intr) warmth physc_ group
warmth     -0.126
physcl_pnsh -0.135 -0.025
group      -0.218 -0.010 -0.019
HDI        -0.925 -0.006  0.008 -0.001
```

11

1.5.3 Julia

In Julia `MixedModels`, the general syntax for a multilevel model of the form described in Equation 1.2 is:

```
fit(MixedModel, @formula(y ~ x + z + (1 + x | group)), data)
```

1.5.3.1 Load The Needed Packages And Load The Data

```julia
using Tables, MixedModels, StatFiles, DataFrames, CategoricalArrays, DataFramesMeta

df = DataFrame(load("simulated_multilevel_data.dta"))
```

1.5.3.2 Graph

```julia
using StatsPlots

@df df scatter(:outcome, :warmth,
               title = "Outcome by Parental Warmth",
               ylabel = "outcome",
               xlabel = "parental warmth")
```

Outcome by Parental Warmth

Figure 1.3: Outcome by Parental Warmth (Julia)

1.5.3.3 Change Country To Categorical

```
@transform!(df, :country = categorical(:country))
```

1.5.3.4 Run The Model

```
m1 = fit(MixedModel, @formula(outcome ~ warmth + physical_punishment +
                  group + HDI +
                  (1 + warmth | country)), df)
```

```
Linear mixed model fit by maximum likelihood
 outcome ~ 1 + warmth + physical_punishment + group + HDI + (1 + warmth | country)
   logLik    -2 logLik     AIC         AICc        BIC
 -9667.9392 19335.8783 19353.8783 19353.9385 19407.9357
```

```
Variance components:
           Column       Variance   Std.Dev.   Corr.
country  (Intercept)    3.2369484 1.7991521
         warmth         0.0001080 0.0103903 +1.00
Residual                36.0187144 6.0015593
 Number of obs: 3000; levels of grouping factors: 30

  Fixed-effects parameters:

                            Coef.   Std. Error        z  Pr(>|z|)

(Intercept)               49.9018    1.43435      34.79   <1e-99
warmth                     0.961545  0.0582135    16.52   <1e-60
physical_punishment       -0.845389  0.0798149   -10.59   <1e-25
group                      1.08524   0.220055      4.93   <1e-06
HDI                        0.0101984 0.0204401     0.50    0.6178
```

2 Cross-Classified Models in Stata, R and Julia

2.1 Introduction

A two level multilevel model imagines that *Level 1* units are nested in *Level 2* units. A three level multilevel model imagines that *Level 1* units are nested in *Level 2* units, which are in turn nested in *Level 3*.

A cross-classified model imagines that the nesting is not hierarchical, but rather that there are two sets of clusters or nestings in which individuals may be nested.

In this data, *events* are nested inside *persons* which are in turn nested in *countries*, since in this data, individuals never change countries. However, the use of a cross-classified framework would allow for a situation in which *persons* moved from country to country, and experienced different *events* in different *countries*.

Below, I describe the use of Stata, R, and Julia to estimate cross-classified models. Because this document is built by Quarto, I describe calling these programs from within a Quarto environment. However, each piece of software could be used individually and separately.

2.2 The Data

The examples below use the `simulated_multilevel_longitudinal_data.dta` file from *Multilevel Thinking*. Here is a direct link to download the data.

Table 2.1: Sample of Simulated Multilevel Longitudinal Data

country	HDI	family	id	group	t	physical_punishment	warmth	outcome
1	69	1	1.1	2	1	2	3	59.18
1	69	1	1.1	2	2	2	2	58.29
1	69	1	1.1	2	3	3	3	60.58
1	69	2	1.2	2	1	4	0	61.54
1	69	2	1.2	2	2	4	0	55.96
1	69	2	1.2	2	3	4	2	56.19

2.3 The Equation

$$\text{outcome}_{ijt} = \beta_0 + \beta_1 t_{ijt} + \beta_2 \text{warmth}_{ijt} + \beta_3 \text{physical punishment}_{ijt} + \beta_4 \text{group}_{ijt} + \beta_5 \text{HDI}_{ijt} + u_{0j} + v_{0i} + e_{ijt}$$

$$(2.1)$$

2.4 Setup

2.4.1 Stata

I need to use the library `Statamarkdown` to call Stata, or I could run Stata on its own

```
library(Statamarkdown)
```

2.4.2 R

In R, I use the library `lme4` to run multilevel models.

```
library(lme4)
```

2.4.3 Julia

I need to call Julia from R.

```
library(JuliaCall)

julia_setup(JULIA_HOME = "/Applications/Julia-1.8.app/Contents/Resources/julia/bin")
```

2.5 Get Data & Run Models

To explain statistical syntax for each software, I consider the more general case of a cross-classified model with dependent variable `y`, independent variables `x` and `z`, clustering variables `country` and `id`.

$$y = \beta_0 + \beta_1 x_{ijt} + \beta_2 z_{ijt} + u_{0j} + v_{0i} + e_{ijt} \qquad (2.2)$$

16

2.5.1 Stata

In Stata `mixed`, the syntax for a multilevel model of the form described in Equation 2.2 is:

```
mixed y x || _all: R.group1 || group2:
```

2.5.1.1 Get The Data

> 💡 Tip For Running Stata From Quarto
>
> Because I am running Stata from inside a Quarto document, and running Stata in multiple chunks, I need to use the `collectcode=TRUE` option in the first Stata chunk. i.e. my Quarto chunk needs to begin with "`{stata, collectcode=TRUE}`
> See Doug Hemken's excellent documentation on `Statamarkdown` here.

```
use simulated_multilevel_longitudinal_data.dta
```

2.5.1.2 Run The Model

```
mixed outcome t warmth physical_punishment group HDI || _all: R.country || id:
```

```
variable t not found
r(111);

end of do-file
r(111);
```

2.5.2 R

In R `lme4`, the general syntax for a multilevel model of the form described in Equation 2.2 is:

```
lmer(y ~ x + z + (1 | group1) + (1 | group2), data = ...)
```

2.5.2.1 Get The Data

```r
library(haven)

df <- read_dta("simulated_multilevel_longitudinal_data.dta")
```

2.5.2.2 Run The Model

```r
fit1 <- lmer(outcome ~ t + warmth + physical_punishment +
               group + HDI +
               (1 | id) +
               (1 | country),
             data = df)

summary(fit1)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: outcome ~ t + warmth + physical_punishment + group + HDI + (1 |
    id) + (1 | country)
   Data: df

REML criterion at convergence: 57088.4

Scaled residuals:
    Min      1Q  Median      3Q     Max
-3.4471 -0.6226  0.0081  0.6153  3.1993

Random effects:
 Groups   Name        Variance Std.Dev.
 id       (Intercept)  8.864   2.977
 country  (Intercept)  3.924   1.981
 Residual             26.008   5.100
Number of obs: 9000, groups:  id, 3000; country, 30

Fixed effects:
                     Estimate Std. Error t value
(Intercept)         49.494782   1.471780  33.629
t                    0.987964   0.065840  15.005
warmth               0.946259   0.038200  24.771
physical_punishment -0.926880   0.049970 -18.549
```

```
group                0.985786   0.153550   6.420
HDI                  0.007543   0.021437   0.352


Correlation of Fixed Effects:
           (Intr) t      warmth physc_ group
t          -0.090
warmth     -0.085  0.008
physcl_pnsh -0.085  0.003 -0.019
group      -0.154  0.000 -0.013 -0.008
HDI        -0.943  0.000 -0.003  0.003  0.000
```

2.5.3 Julia

In Julia `MixedModels`, the general syntax for a multilevel model of the form described in Equation 2.2 is:

```
fit(MixedModel, @formula(y ~ x + z + (1 | group1) + (1 | group2)), data)
```

2.5.3.1 Load The Needed Packages And Load The Data

```julia
using Tables, MixedModels, StatFiles, DataFrames, CategoricalArrays, DataFramesMeta

df = DataFrame(load("simulated_multilevel_longitudinal_data.dta"))
```

2.5.3.2 Change Country To Categorical

```julia
@transform!(df, :country = categorical(:country))
```

2.5.3.3 Run The Model

```julia
m1 = fit(MixedModel, @formula(outcome ~ t + warmth + physical_punishment +
                              group + HDI +
                              (1 | id) +
                              (1 | country)), df)
```

19

```
Linear mixed model fit by maximum likelihood
 outcome ~ 1 + t + warmth + physical_punishment + group + HDI + (1 | id) + (1 | country)
     logLik    -2 logLik        AIC          AICc         BIC
 -28533.9968   57067.9935   57085.9935   57086.0136   57149.9384


Variance components:
           Column    Variance Std.Dev.
id       (Intercept)   8.85264 2.97534
country  (Intercept)   3.65030 1.91058
Residual               26.00093 5.09911
 Number of obs: 9000; levels of grouping factors: 3000, 30

  Fixed-effects parameters:

                          Coef.   Std. Error       z  Pr(>|z|)

(Intercept)              49.4945      1.42422     34.75    <1e-99
t                         0.987965    0.0658315   15.01    <1e-50
warmth                    0.946255    0.0381869   24.78    <1e-99
physical_punishment      -0.926774    0.0499549  -18.55    <1e-76
group                     0.985819    0.153487     6.42    <1e-09
HDI                       0.00754357  0.0207101    0.36    0.7157
```